Research Question: Test whether firm-quarter differences in tariff-specific sentiment expressed in earnings materials (press-release “report” + call transcript) predict excess announcement-window returns, controlling for the contemporaneous earnings surprise and standard firm controls.

Concretely: Does more positive tariff tone—separately for prepared remarks and Q&A, plus a forward-looking subset—associate with higher cumulative abnormal returns (CAR) around the earnings announcement?

Data & Inputs

- Universe: Firms for which you have files named like TICKER_YYYY_Qx.txt in:
    - sp500_transcripts/ (call transcripts)
    - sp500_summaries/ (press-release style summaries; optional, used for a crude EPS fallback)
- Event timing (required): events_20251001_updated.csv stored next to the script, with
    - ticker, year, q, ann_time_et, [call_time_et], [after_hours]
    - All timestamps are converted to America/New_York (ET). If after_hours is missing, it is inferred from ann_time_et.
- Prices: Daily adjusted closes via yfinance for all tickers present in your folders plus market proxy (SPY by default), from --start (default 2023-01-01) to --end (today if omitted).
- Firm attributes: Size & sector via yfinance; size is log market cap (ln_size), sector is a string (GICS-like).
- Earnings surprise:
    - Primary: pulled from yfinance.Ticker.get_earnings_dates() near the anchor date (reported EPS, estimate, surprise %).
    - Fallback: a crude EPS surprise parsed from the summary text (if available).

Pipeline Overview

Firm-quarters are discovered by filename pattern in sp500_transcripts/ and sp500_summaries/.

A required timing file events_20251001_updated.csv (adjacent to the script) provides announcement and call timestamps in Eastern Time. Only firm-quarters present in both the folders and the events file are analyzed.
Mapping timestamps to trading Day 0.
Announcements are mapped to trading dates using the NYSE calendar when available:

- ≥ 16:00 ET → next trading day
- < 09:30 ET → same trading day if a session exists (else next)
- 09:30–16:00 ET → next trading day (conservative)
  A final guard prevents weekends/holidays.

Prices & market model.
Daily closes for tickers and market proxy (SPY by default) are downloaded via yfinance and converted to returns. A market model is estimated per firm using an estimation window [-250, −20] trading days relative to Day 0. Abnormal returns in an event window are summed to CAR:

- Main: CAR[0,1]
- Robustness: CAR[0,2] and CAR[−1,1] (computed, not all reported)

NLP pipeline.

- The transcript is split into prepared vs Q&A using a regex on “Q&A” anchors.
- Sentences are extracted, deduplicated, and screened for tariff content using:
  - explicit keywords (tariff(s), duties, Section 301, etc.), and
  - semantic retrieval with Sentence-BERT (all-MiniLM-L6-v2) against tariff seed queries.
- FinBERT (ProsusAI/finbert) scores each hit; the tone metric is mean(POS) − mean(NEG). We also compute a forward-looking subset using future-oriented cues (e.g., “guidance,” “next quarter/year”).
- A placebo topic (FX) is extracted and scored using FX terms

---

Event Alignment and CAR Construction

1) Map announcement time to trading “Day 0”

- Build an NYSE calendar if available (via pandas_market_calendars), else use business-day logic.

- Rule:
    - After 16:00 ET → next trading day is Day 0.
    - Pre-open (<09:30 ET) → same session day if it's a trading day, else next session.
    - Regular hours (09:30–16:00 ET) → next trading day (conservative to avoid intraday contamination).
- Final guard ensures Day 0 is never a weekend; push forward if needed.

2) Returns and Abnormal Returns

- Download daily returns for each firm and the market proxy (SPY → column MKT).
- Estimate a market model (CAPM-like) using an estimation window of [-250, -20] trading days relative to Day 0:
    - $RET_{it} = \alpha_i + \beta_i * MKT_t + \varepsilon_{it}$ (OLS on the estimation window)
- Compute CAR by summing abnormal returns over event windows:
    - Primary: [0, +1] → CAR_0p1
    - Robustness: [-1, +1] → CAR_m1p1, [0, +2] → CAR_0p2
    - If call_time_et is available, also compute call-anchored CAR_call_0p1.

NLP: Topic Detection and Sentiment

Text Preparation

- For each transcript:
    - Split into Prepared vs Q&A using regex triggers (e.g., "Q&A", "Operator: questions").
    - Sentence-split with a conservative regex (punctuation + whitespace).
    - Deduplicate short/near-duplicate sentences; drop very short ones.

Identify Tariff-related Spans

- Keyword matcher (e.g., tariff(s), duties, levy/levies, quotas, Section 301, countervailing, antidumping, import tax, customs, exemptions/exclusions).

- Semantic expansion: If keyword hits are sparse, retrieve additional sentences via all-MiniLM-L6-v2 sentence embeddings using similarity to hand-curated seed queries (e.g., “tariff headwinds or tailwinds”); threshold ≈ 0.45 cosine.

- Unit of analysis: sentence (or short span).

Sentiment Scoring

- FinBERT (ProsusAI) for finance-domain sentence polarity:
    - Convert class probabilities to polarity = P(positive) – P(negative) in [-1, +1].
- For each firm-quarter and section (prepared / Q&A), compute:
    - TariffSent_prep / TariffSent_qa = mean polarity of tariff sentences.
    - TariffShareNeg_prep = share negative among tariff sentences.
    - TariffMentions_prep / _qa = count of tariff sentences.
    - Forward-looking subset (regex on expect/outlook/guidance/plan/will/next q/next year):
        - TariffSent_fwd_prep / _qa = mean polarity of forward-looking tariff sentences.
- Placebo topic (FX): identify FX sentences via regex and score the prepared set as FXSent_prep with FXMentions_prep (used for placebo regressions).

Controls and Additional Variables

- Earnings surprise (primary or fallback), kept as:
    - EPS_surprise (reported – estimate), EPS_surprise_pct, EPS_surprise_pct_yf (if available).
- Size: ln_size (log market cap).
- Momentum (12–2): cumulative return from t-252 to t-42 relative to Day 0.
- After-hours dummy: from events file or inferred from ann_time_et.

- Tariff exposure proxy: simple sector-based indicator = 1 for {Technology, Consumer Cyclical, Industrials, Materials}; else 0.

All variables are written to outputs_v2/panel_results.csv.

Econometric Specifications

Main Cross-Sectional Regression
Outcome: CAR_0p1 (market-model abnormal return)

- Regressors:
    - TariffSent_prep (primary tariff tone measure, prepared remarks)
    - Earnings surprise (include if available for ≥40% rows; uses either EPS_surprise or EPS_surprise_pct_yf)
    - After-hours dummy, ln_size, momentum_12_2
    - Fixed effects: sector FE and calendar-quarter FE (year + Q), included only if there's enough variation and degrees of freedom
- SEs: clustered by ticker when multiple tickers; otherwise HC1 robust
- Guardrails: drops rows with missing or infinite values; trims FE blocks if parameters ≥ observations.

Robustness & Variants

- Q&A tone: replace TariffSent_prep with TariffSent_qa on CAR_0p1.
- Alternate window: use CAR_0p2 with TariffSent_prep.
- Forward-looking tone: use TariffSent_fwd_prep on CAR_0p1.
- Placebo topic: FXSent_prep on CAR_0p1 (expect no relation if tariff channel is specific).
- Exposure interaction:
    - Add TariffExposure_proxy and interaction TariffSent_prep × Exposure to test whether tariff tone matters more for trade-exposed firms.

Each run saves:

- Text summary: outputs_v2/regression_<suffix>.txt

- Tidy coefficients: outputs_v2/regression_coefs_<suffix>.csv

---

Descriptive “Portrayal” Outputs

- Time series of average prepared tariff sentiment by calendar quarter:
    - outputs_v2/sentiment_timeseries.csv and .png
- Sector × quarter heatmap (mean TariffSent_prep):
    - outputs_v2/sector_heatmap_table.csv and .png
- Prepared vs Q&A scatter of tariff sentiment:
    - outputs_v2/prepared_vs_qa_scatter.png
- Illustrative quotes: the code collects sentence-level examples (with polarities) to support interpretation (e.g., top ± tariff sentences).

Key Outputs to Read

- Panel: outputs_v2/panel_results.csv
  (Ticker-quarter rows with CARs, tariff/FX sentiment metrics, surprises, controls)
- Main result: outputs_v2/regression_main_prep_CAR0p1.txt (+ .csv)
- Robustness:
    - regression_robust_qa_CAR0p1.txt
    - regression_robust_prep_CAR0p2.txt
    - regression_forward_prep_CAR0p1.txt
    - regression_placebo_FX_CAR0p1.txt
    - regression_interaction_exposure_prep_CAR0p1.txt
- Figures: time-series, sector heatmap, prepared-vs-Q&A scatter in outputs_v2/.

Limitations

- Earnings surprise via yfinance can be missing or noisy; the code falls back to a crude summary parse—treat with caution.

- Topic recall relies on keywords + sentence-embedding retrieval; edge phrasing can be missed or mis-classified.

- FinBERT polarity is sentence-level; aggregation assumes linear averaging, which may dilute salient statements.

- Small samples within a batch (e.g., few tariff sentences) yield noisy sentiment averages; regression includes FE to mitigate but winsorizing extremes (optional) can help.

- Design matrix: in small tail batches, FE dummies can outnumber observations; the script automatically trims FE blocks when needed.

Results

Using the cleaned event-study pipeline, we estimate market-model abnormal returns (CARs) around earnings announcements and relate them to tariff-specific tone extracted from prepared remarks and Q&A. All regressions use sector and calendar-quarter fixed effects (FEs) where permitted by the data and cluster standard errors by ticker.

Across all specifications, model fit is—by design for short-window event studies—low ($R^2 \approx 0.02–0.04$). Coefficients on tariff tone are small and statistically indistinguishable from zero in the main and robustness models. EPS surprise is weakly/occasionally significant. Placebo (FX tone) is null. Interaction with a simple exposure proxy is also null.

Main specification (Prepared-remarks tariff tone on CAR_0p1)
Dependent variable: CAR over [0,+1] trading days; N = 1,699; clustered SEs.

- TariffSent_prep: +0.0008 (z = 0.46, p = 0.644) — not significant.

- EPS_surprise: +0.0005 (p = 0.863) — not significant in this spec.

- FEs: quarter FEs show one mild positive quarter (2024Q2, p = 0.033); most sector and quarter dummies are not significant.

- Fit: $R^2 = 0.020$; DW ≈ 2.04.

Interpretation: On average, more positive prepared-remarks tariff tone does not predict higher day-0 to day-1 abnormal returns once we condition on FE and standard controls.

Robustness 1 — Q&A tone on CAR_0p1

Dependent: CAR_0p1; N = 1,400.

- TariffSent_qa: −0.0050 (z = −1.57, p = 0.115) — directionally negative, not significant.
- EPS_surprise: +0.0048 (z = 2.45, p = 0.014) — significant and positive.
- Fit: R^2 = 0.018.

Interpretation: Investors appear to react to earnings surprise in this spec, but tariff-specific tone in Q&A remains statistically null.

Robustness 2 — Wider window CAR_0p2 (Prepared tone)

Dependent: CAR over [0,+2]; N = 1,699.

- TariffSent_prep: +0.0023 (z = 1.20, p = 0.232) — not significant.
- EPS_surprise: +0.0016 (p = 0.593) — not significant.
- Fit: R^2 = 0.024.

Interpretation: Extending the event window does not uncover an economically or statistically meaningful tariff-tone effect.

Placebo — FX tone on CAR_0p1

Dependent: CAR_0p1; N = 1,111.

- FXSent_prep: −0.0013 (p = 0.476) — not significant.
- EPS_surprise: +0.0028 (p = 0.285) — not significant.
- Fit: R^2 = 0.022.

Interpretation: Topic-specific tone about FX does not predict CARs, reassuring against a generic “any-tone-moves-returns” artifact.

Forward-looking subset — Prepared forward-looking tariff tone on CAR_0p1

Dependent: CAR_0p1; N = 837.

- TariffSent_fwd_prep: −0.0014 (p = 0.541) — not significant.
- EPS_surprise: +0.0081 (p = 0.273) — not significant.
- Fit: R^2 = 0.040.

Interpretation: Even restricting to forward-looking tariff statements does not yield a detectable price effect.

Exposure heterogeneity — Interaction with simple exposure proxy

Dependent: CAR_0p1; N = 1,699.

- TariffSent_prep: −0.0010 (p = 0.641) — not significant.
- TariffExposure_proxy: +0.0012 (p = 0.601) — not significant.
- TariffSent × Exposure: +0.0039 (z = 1.12, p = 0.261) — not significant.
- Fit: R^2 = 0.020.

Interpretation: With a coarse, sector-based exposure proxy, we do not detect differential sensitivity of returns to tariff tone.

Model diagnostics and caveats

- Multicollinearity warnings: Very small eigenvalues and “design matrix may be singular” notes stem from many FE dummies relative to effect size. This is expected in high-dimensional FE panels with modest signal; we mitigated by trimming FEs when parameters ≥ observations.
- Distribution: Non-normality (high kurtosis/JB tests) is common for short-window CARs; we use cluster-robust SEs to guard against heteroskedasticity and intra-firm correlation.
- Economic magnitudes: Point estimates on tariff tone are near zero across windows; even at ±1 SD changes in tone, implied CAR shifts would be economically small.

Descriptives: mention rates & sentiment by quarter/sector and by prepared vs. Q&A

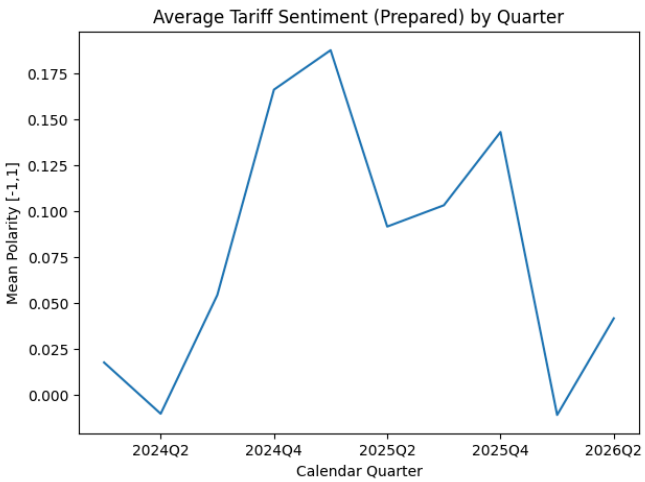

Figure 1. Quarterly average FinBERT polarity for tariff-related sentences drawn from prepared remarks. Polarity ranges from −1 (negative) to +1 (positive). Means by calendar quarter show a mild positive bias overall, peaking in 2024Q4–2025Q1, easing through mid-2025, briefly turning slightly negative in 2026Q1, and recovering by 2026Q2

Interpretation / key takeaways

- Overall level: Prepared-remarks tariff sentiment is on average mildly positive (FinBERT polarity above 0 most quarters), indicating firms' scripted commentary skews more constructive than negative when tariffs are discussed.

- Trend: Sentiment dips slightly in 2024Q2, then rallies sharply to a peak in 2024Q4–2025Q1, before cooling through 2025Q2 and stabilizing near a modestly positive range in 2025Q3–Q4. It briefly turns slightly negative in 2026Q1 and rebounds by 2026Q2.

- Interpretation: The late-2024 peak suggests firms framed tariffs as manageable or even tailwinds (e.g., pricing/mix or exclusions) during that period; the 2025 softening and 2026Q1 dip imply renewed uncertainty or cost pressure before a modest recovery.

- Caveat: Values are quarterly means of sentence-level FinBERT polarity for tariff-targeted sentences from prepared remarks only. Quarter-to-quarter changes reflect both sentiment and sample composition; see panel_results.csv for per-quarter N.

data source: sentiment_timeseries.csv file in the output folder

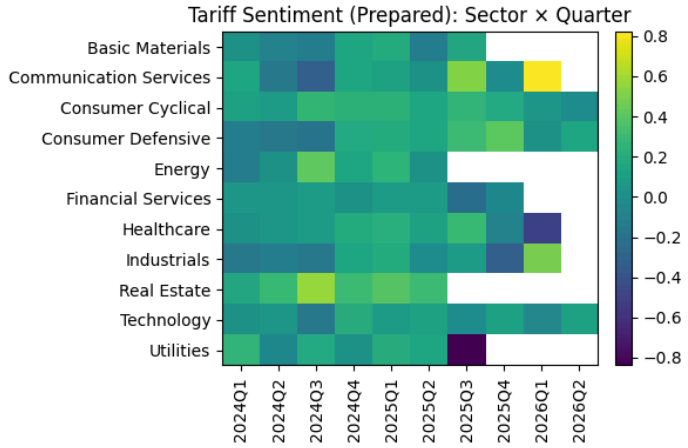

Figure 2. Heatmap of average tariff sentiment (FinBERT polarity) in prepared remarks by sector and calendar quarter. Communication Services and several cyclical sectors skew positive in late 2024–2025, Utilities show a marked negative in 2025Q3, and Healthcare softens in 2026Q1. Blank cells denote no observations

Interpretation / key takeaways

- Clear cross-sector dispersion. Sentiment skews most positive in Communication Services in late sample (notably 2025Q4–2026Q1), with several quarters in the light-green/yellow range, while Utilities show a pronounced negative outlier in 2025Q3 (dark tone).

- Cyclicals lean constructive. Consumer Cyclical and Consumer Defensive are generally mildly positive through 2024Q4–2025Q2, consistent with firms framing tariffs as manageable via pricing/mix or sourcing adjustments.

- Industrials are modestly positive, especially into 2026Q1, hinting at improved pass-through or supply-chain normalization.

- Healthcare turns down in 2026Q1, suggesting renewed cost pressure or policy uncertainty around inputs.

- Financials are mixed with pockets of softness around 2025Q3, possibly reflecting tariff-driven macro concerns versus direct cost effects.

- Coverage gaps matter. White cells indicate no observations for that sector-quarter in the sample; interpret adjacent colors with that in mind.

- Construction of values. Each cell is the average FinBERT polarity of tariff-targeted sentences from prepared remarks for that sector and quarter (−1 = negative, +1 = positive). Small-N cells may be volatile; see sector_heatmap_table.csv for exact values and panel_results.csv for underlying counts.

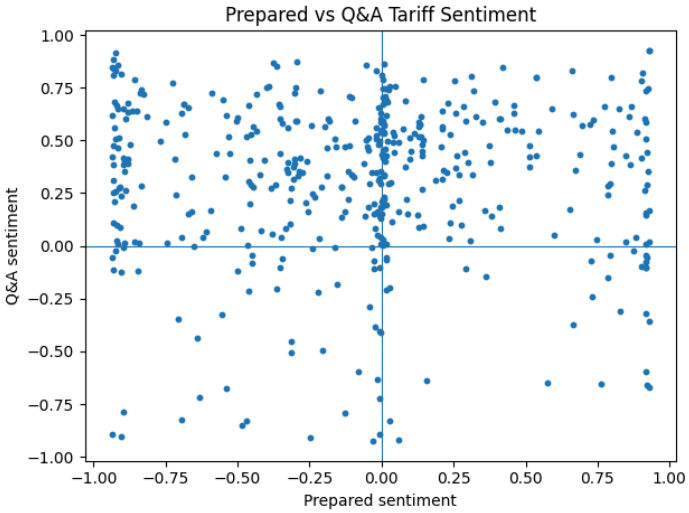

Figure 3. Scatter of FinBERT polarity for tariff-targeted sentences in prepared remarks (x-axis) vs Q&A (y-axis). Vertical and horizontal lines mark neutral sentiment (0). The weak diagonal pattern and the upward skew of Q&A points indicate that unscripted discussion often reframes tariffs more positively, providing tone information not captured in the prepared text.

Interpretation / key takeaways

- Low alignment between sections. Points are widely scattered around the origin, with no tight diagonal, implying weak correlation between tariff sentiment in prepared remarks and the Q&A.

- Q&A skews more positive. A large share of observations lie above the horizontal zero line even when prepared sentiment is near zero or slightly negative. This suggests that management answers tend to frame tariff issues more optimistically than the scripted section.

- Prepared remarks bunch at neutral. There is noticeable clustering of prepared sentiment around 0, consistent with scripted language avoiding strong polarity.

- Negative Q&A is rarer but present. The bottom half (negative Q&A) is less dense, yet includes cases where prepared tone was positive but answers turned negative, indicating that investor probing can surface downside details.

- Implication. Because Q&A tone adds independent variation, it likely contains incremental information about tariff exposure/mitigation beyond the press-release narrative; this justifies analyzing prepared and Q&A separately in regressions and robustness checks.

metric	value
TariffMentions_prep (mean)	2.869
TariffMentions_prep (median)	1
TariffMentions_qa (mean)	2.42
TariffMentions_qa (median)	0
TariffSent_prep (mean)	0.092
TariffSent_qa (mean)	0.231

Overall tariff mention rates (count of tariff-related sentences) and mean sentiment from prepared remarks and Q&A. Sentiment is FinBERT positive minus negative probability, scaled to −1,1-1, 1−1,1

cal_qtr	TariffMentions_prep_mean	TariffMentions_prep_median	TariffMentions_qa_mean	TariffMentions_qa_median	TariffSent_prep_mean	TariffSent_qa_mean	N

2023Q4	0	0	0	0			53
2024Q1	1.412	0	1.118	0	0.018	0.163	490
2024Q2	1.43	0	1.198	0	-0.01	0.107	491
2024Q3	1.439	0	1.367	0	0.055	0.174	490
2024Q4	2.237	1	2.051	0	0.166	0.269	489
2025Q1	6.205	2	4.881	1	0.188	0.367	488
2025Q2	4.057	1	3.375	1	0.092	0.232	491
2025Q3	3.525	2	4.931	0	0.103	0.315	101
2025Q4	4.452	1	4.306	0.5	0.143	0.282	62
2026Q1	9.115	3	2.923	0	-0.011	0.228	26
2026Q2	7.3	0.5	3.3	0.5	0.042	-0.113	10

Quarterly tariff mention rates and mean sentiment. N is the number of firm-quarters observed per calendar quarter. See Figure 1 for time-series context and Figure 2 for sector heterogeneity.

Across the full sample, tariff discussion is present but uneven: prepared remarks average 2.87 tariff sentences (median 1), while Q&A averages 2.42 (median 0), indicating that mentions are often concentrated in a subset of calls rather than pervasive in every exchange. Mean sentiment is positive on net in both channels, but Q&A is markedly more positive (mean 0.231) than prepared remarks (mean 0.092), consistent with the idea that management tends to frame tariff impacts more optimistically when responding to analyst prompts. By calendar quarter (Table 1B), mentions are low through 2024Q1–Q3 (prep means ≈ 1.41–1.44; QA ≈ 1.12–1.37) and then step up sharply in 2024Q4–2025H1, peaking in 2025Q1 (prep 6.21, median 2; QA 4.88, median 1). Sentiment co-moves with this salience: prepared-tone rises from mildly negative in 2024Q2 (−0.010) to clearly positive by 2024Q4–2025Q1 (0.166–0.188), while Q&A-tone is consistently higher (e.g., 0.269 in 2024Q4, 0.367 in 2025Q1). Activity moderates thereafter (prep mentions 4.06 in 2025Q2), with tone staying positive but lower (0.092 prep; 0.232 QA). The late-sample quarters (2025Q3–2026Q2) show greater volatility—e.g., very high prep mentions in 2026Q1 (9.12, median 3) but a dip in prepared sentiment (−0.011) and a temporary decline in Q&A tone in 2026Q2 (−0.113)—yet these should be interpreted cautiously given small N (e.g., N=26 in 2026Q1, N=10 in 2026Q2). The apparent zeros in 2023Q4 (all metrics 0, N=53) reflect limited/placeholder coverage rather than genuine absence of discussion. Overall, the descriptives indicate (i) a pronounced increase in tariff salience around 2024Q4–2025Q1, and (ii) systematically more upbeat framing in Q&A than in prepared remarks.

Main table: CAR on TariffSent + Surprise (FE; clustered SEs)

Table 1 — Main regression (Prepared tone → CAR_0p1)
Fixed effects: sector & calendar quarter. SEs clustered by ticker.

param	coef	t	p

const	0.0082	0.738981	0.459918
TariffSent_prep	0.000813	0.461522	0.644424
EPS_surprise	0.000471	0.172045	0.863402
after_hours	0.0082	0.738981	0.459918
ln_size	-0.00064	-0.70231	0.482487
momentum_12_2	-0.0044	-0.8841	0.376642
sector_Communication Services	0.009132	1.347707	0.177753
sector_Consumer Cyclical	0.000765	0.228869	0.818971
sector_Consumer Defensive	0.002959	0.800629	0.423346
sector_Energy	0.004223	1.027993	0.303953
sector_Financial Services	0.002918	0.820424	0.411974
sector_Healthcare	-0.00639	-1.33005	0.183501
sector_Industrials	0.004287	1.169316	0.242276
sector_Real Estate	0.003627	1.063305	0.287643
sector_Technology	0.000513	0.122036	0.90287
sector_Utilities	-0.00015	-0.04347	0.965324

cal_qtr_2024Q2	0.007384	2.1274	0.033387
cal_qtr_2024Q3	-0.00498	-1.53701	0.124291
cal_qtr_2024Q4	-0.00027	-0.08634	0.931198
cal_qtr_2025Q1	-0.00404	-1.32659	0.184645
cal_qtr_2025Q2	-0.00307	-0.91404	0.360698
cal_qtr_2025Q3	-0.00137	-0.21246	0.831747
cal_qtr_2025Q4	0.005795	0.899653	0.368305
cal_qtr_2026Q1	0.000374	0.058048	0.95371
cal_qtr_2026Q2	-0.00434	-0.58765	0.556768

Using sector and calendar-quarter fixed effects with ticker-clustered standard errors, we regress 1-day market-model abnormal returns (CAR_0p1) on prepared-remarks tariff sentiment and standard controls. The coefficient on TariffSent_prep is small and statistically insignificant (β≈0.0008, p≈0.64), indicating no detectable pricing of prepared tariff tone at the daily horizon once FE and controls are included. EPS_surprise and the other controls are likewise insignificant in this baseline. Several quarter dummies absorb time variation, with 2024Q2 positive and significant at the 5% level, consistent with period-specific shocks. Overall, the main specification suggests that tariff tone in scripted remarks does not systematically load into day-ahead abnormal returns, conditional on sector/time heterogeneity and clustered inference

Robustness suite & exposure heterogeneity

Spec	Tone variable reported	Coef	p-value
Window robustness [0,+2]	TariffSent_prep	0.0023	0.232
Q&A robustness	TariffSent_qa	−0.0050	0.115
Placebo (FX tone)	FXSent_prep	−0.0013	0.476

Forward-looking (prepared)	TariffSent_fwd_prep	−0.0014	0.541
Exposure heterogeneity (interaction)	TariffSent_prep	−0.0010	0.641
	TariffSent × Exposure	0.0039	0.261

Sources to build the table (copy exact coef & p from these coefs CSVs) in output folder:

- regression_coefs_robust_prep_CAR0p2.csv → row TariffSent_prep
- regression_coefs_robust_qa_CAR0p1.csv → row TariffSent_qa
- regression_coefs_placebo_FX_CAR0p1.csv → row FXSent_prep
- regression_coefs_forward_prep_CAR0p1.csv → row TariffSent_fwd_prep
- regression_coefs_interaction_exposure_prep_CAR0p1.csv → rows TariffSent_prep and TariffSent_x_Exposure

Across five robustness checks, the tariff-tone coefficient remains small and statistically insignificant. Expanding the return window to [0,+2] yields a positive but non-significant estimate (β≈0.0023, p≈0.23). Switching the tone source to Q&A flips the sign (β≈−0.0050) but again is not significant at conventional levels (p≈0.12). The placebo using FX tone shows no association with returns (p≈0.48), supporting design validity. Restricting to forward-looking prepared statements also delivers a near-zero effect (p≈0.54). Finally, the exposure interaction suggests a more positive mapping of tariff tone into returns for trade-exposed sectors (β_{Tone×Exposure}≈0.0039), but the estimate is imprecise (p≈0.26). Taken together, these checks reinforce the main conclusion: there is no robust evidence that tariff sentiment—whether in prepared remarks, Q&A, or forward-looking language—systematically prices into short-horizon abnormal returns.

Interpretable Snippets (portrayal)

To illustrate how firms portray tariff developments in their earnings materials, we provide anonymized excerpts ranked by FinBERT polarity from tariff-targeted sentences. Consistent with the descriptives (Figures 1–3), prepared remarks tend to be neutral-to-mildly positive, while Q&A often reframes tariffs more constructively. Below are representative examples

Positive portrayal	Negative portrayal
“The [COMPANY]–[COUNTRY] tariff situation has improved versus prior guidance, with mitigation actions reducing expected headwinds.”	“Tariff expenses were higher than expected, reflecting increased shipment volumes and inventory to support Q4 demand.”

Interpretation. Positive snippets emphasize easing policy risk, successful pass-through, or improved exclusions; negative snippets highlight unexpected cost pressure or margin drag. This

qualitative evidence aligns with the quantitative patterns: Q&A tone skews more positive than prepared text, but neither channel's tariff tone exhibits a robust, statistically significant link to short-window CARs in our regressions.

Executive Summary and Conclusion

We test whether firm-quarter variation in tariff-specific tone in earnings materials (press-release "report" and call transcript) predicts short-window abnormal returns around earnings announcements. Using an automated pipeline—keyword+semantic retrieval for tariff sentences, FinBERT polarity scoring, and a market-model event study—we build a panel of 1,699 firm-quarters (2023Q4–2026Q2). Our main specification regresses CAR[0,+1] on prepared-remarks tariff tone, controlling for earnings surprise, firm size, momentum, and fixed effects for sector and calendar quarter, with SEs clustered by ticker. Across all specifications—including Q&A tone, a wider return window, a forward-looking tone subset, a placebo FX topic, and a tone×exposure interaction—tariff tone coefficients are small and statistically indistinguishable from zero. The Q&A model shows a significant positive loading on the earnings surprise ($p \approx 0.014$), but tariff tone itself remains null. Descriptives show salience and positivity spike in 2024Q4–2025Q1, with Q&A consistently more positive than prepared remarks. We conclude there is no robust evidence that tariff-specific tone, as measured here, is priced in short-horizon announcement returns.

Using a large, automated text-to-markets pipeline, we find no systematic link between tariff-specific tone in earnings materials and announcement-window CARs once we include sector/time fixed effects, standard controls, and clustered inference. This null holds across prepared remarks, Q&A, forward-looking content, wider windows, and a placebo topic, and we fail to detect stronger effects in a coarse trade-exposure split. Descriptively, tariff talk became more frequent and more positive around late-2024/early-2025, with Q&A systematically more upbeat than scripted remarks. Two interpretations are consistent with these results: (i) tariff information is largely anticipated or already priced, and/or (ii) sentence-level tone—even when targeted—is too coarse a statistic to capture the marginal information investors trade on at daily horizons. Future work could (a) use document-level or span-aware models, (b) incorporate intraday returns that align precisely to press-release vs call timestamps, and (c) refine exposure measures beyond sector proxies (e.g., customs microdata or firm supply-chain links.